

AUTOTUNING: A DESIGN OF EXPERIMENTS APPROACH

Pedro Bruel

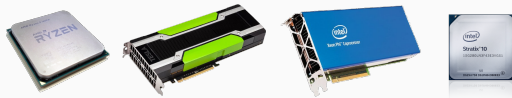
phrb@ime.usp.br

Journée au Vert POLARIS, March 2018

1. Autotuning
2. Applying Design of Experiments to Autotuning
3. Perspectives

AUTOTUNING: OPTIMIZING PROGRAM CONFIGURATION

Architectures for High Performance Computing



How to write **efficient code** for each of these?

Autotuning

The process of **automatically finding** a **configuration** of a program that optimizes an **objective**

Configurations

- Program configuration
 - Algorithm, block size, . . .
- Source code transformation
 - Loop unrolling, tiling, rotation, . . .
- Compiler configuration
 - -O2, vectorization, . . .
- . . .


Objectives

- Execution time
- Memory & power consumption
- . . .

Search Spaces

Represent the **effect** of all possible **configurations** on the **objectives**

Can be difficult to explore, with multiple **local optima** and **undefined regions**



Hölder Table function

Issue 1: Exponential Growth

Simple factors can generate large spaces:

- 30 *boolean* factors
- 2^{30} combinations

Issue 2: Geometry

- Discrete or continuous factors
- “Smoothness”
- Interactions between factors

Issue 3: Measurement Time

Time to compile:

- Benchmark GPU applications: 1~10s
- Benchmark FPGA applications: 1~10min
- Industrial FPGA applications: 1~10h

AUTOTUNING: MULTIPLE APPROACHES

Popular Approaches

- Exhaustive
- Meta-Heuristics
- Machine Learning

System	Domain	Approach
ATLAS	Dense Linear Algebra	Exhaustive
INSIEME	Compiler	Genetic Algorithm
Active Harmony	Runtime	Nelder-Mead
ParamILS	Domain-Agnostic	Stochastic Local Search
OPAL	Domain-Agnostic	Direct Search
OpenTuner	Domain-Agnostic	Ensemble
MILEPOST GCC	Compiler	Machine Learning
Apollo	GPU kernels	Decision Trees

Main Issues

- Optimized function is a **black-box**:
 - Learn nothing about the search space
 - Can't explain why optimizations work
- These approaches **assume**:
 - A large number of function evaluations
 - Search space “smoothness”
 - Good solutions are reachable

APPLYING DESIGN OF EXPERIMENTS TO AUTOTUNING

Our Approach

Using **efficient experimental designs** to overcome issues related to **exponential growth**, **geometry**, and **measurement time**

Design Requirements

- Support a large number of factors (**Exponential Growth**)
- Support continuous and discrete factors (**Geometry**)
- Minimize function evaluations (**Measurement Time**)

Main Design Candidates

Screening Designs:

- Assume **interactions are negligible**
- Estimate **main effects**
- Aim to **minimize runs**

Mixed-Level Designs:

- Factors have **different number of levels**
- Many **optimality criteria**

SCREENING DESIGNS

A Plackett-Burman **screening design** for 7
2-level factors:

Run	A	B	C	D	E	F	G
1	1	-1	1	-1	-1	1	1
2	1	1	1	-1	1	-1	-1
3	-1	1	-1	-1	1	1	1
4	-1	1	1	1	-1	1	-1
5	1	-1	-1	1	1	1	-1
6	1	1	-1	1	-1	-1	1
7	-1	-1	1	1	1	-1	1
8	-1	-1	-1	-1	-1	-1	-1

Screening Designs

Plackett-Burman designs for 2-level factors:

- Orthogonal arrays of strength 2
- Estimate the main effects of n factors with $n + 1$ runs


Construction:

- For $n + 1$ multiple of 4
- Identical to a fractional factorial design if $n + 1$ is a power of two

LOOKING AT DATA: CUDA COMPILER FLAGS

CUDA Compiler Flags

- Rodinia Benchmark
- 16 factors, few with multiple levels
- 10^6 combinations
- 1~10s to measure
- Screening Experiment:
 - 16 “2-level” factors
 - 4 “dummy” factors



MIXED-LEVEL DESIGNS

A multi-level design for 1 2-level factor and 3 3-level factors:

Run	A	B	C	D
1	1	1	1	3
2	1	1	2	1
3	1	1	3	2
4	1	2	1	2
5	1	2	2	3
6	1	2	3	1
7	1	3	1	1
8	1	3	2	2
9	1	3	3	3
10	2	1	1	1
11	2	1	2	2
12	2	1	3	3
13	2	2	1	3
14	2	2	2	1
15	2	2	3	2
16	2	3	1	2
17	2	3	2	3
18	2	3	3	1

Mixed-Level Designs

Strategy 1: Contractive Replacement

- Find specific sets of k -level columns of a design
- Contract the set into a new factor with more levels
- Maintain orthogonality of the design

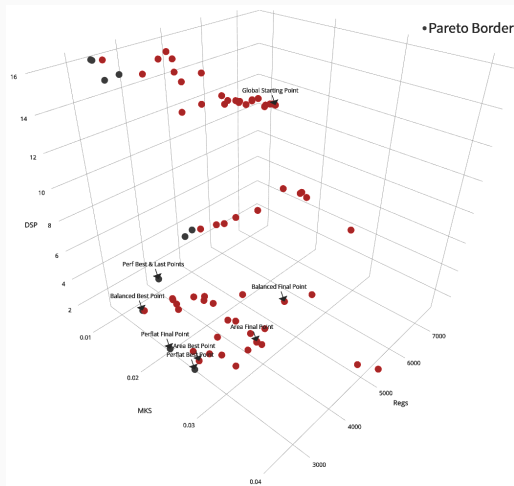
Strategy 2: Direct Construction

Directly generate small mixed-level designs by solving Mixed Integer Programming problems

LOOKING AT DATA: FPGA COMPILER PARAMETERS

FPGA Compiler Parameters

- CHStone Benchmark
- 141 factors, most with multiple levels
- 10^{128} combinations
- 1~10min to measure
- Multiple objectives
- Search with Meta-Heuristics:
 - Unstructured data difficult analysis
 - We are working on obtaining more data



Perspectives

- **Short term:**
 - Study **small, balanced, orthogonal multi-level** designs for **large numbers of factors**
 - Iteratively **drop least significant factors** with **user input**
- **Long term:**
 - Use such designs to **autotune industrial-level FPGA applications**
 - Provide an **autotuning shared library** to applications

Takeaway

Target Scenario: **FPGA Compiler Parameters**

- **Large search space**
- Factors with **multiple levels**
- **Large measurement time**

Our Approach

Using **efficient experimental designs** to overcome issues related to **exponential growth, geometry**, and **measurement time**

Main Design Candidates

Screening & Mixed-Level designs

AUTOTUNING: A DESIGN OF EXPERIMENTS APPROACH

Pedro Bruel

phrb@ime.usp.br

Journée au Vert POLARIS, March 2018